

Introduction to OpenCV

By Noriaki Kono

Computer Vision (CV)

Why CV?

Pros:

- ▶ Needs very few images to “learn”
- ▶ All results are expected and understood

Cons:

- ▶ Generalized solution can complicate the model
- ▶ Rigid and inflexible model
 - ▶ Requires new solutions for new situations
- ▶ Can be very difficult to achieve optimal results

Image Formation

What is depth?



- ▶ Images are 2D, implying no depth
- ▶ Human brains assume depth
 - ▶ Object overlap
 - ▶ Shadows

Image Formation

What are edges?



- ▶ Two very basic types of edges:
 - ▶ Hard Edges - Edges which transition quickly
 - ▶ Soft Edges - Edges which transition slowly
- ▶ Derivatives help measure edges

Image Formation

What are edges?



Hard Edge Soft Edge

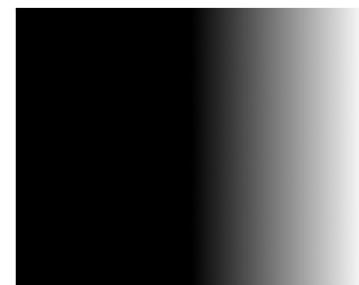


Image Formation

What are edges?



How many edges?



Image Formation

What is a pixel?

- ▶ Measures Intensity
 - ▶ Information carrier
- ▶ A byte translates to 256 values (0-255)
- ▶ 0 - represents black
- ▶ 255 - represents white

Image Formation

What is a pixel?

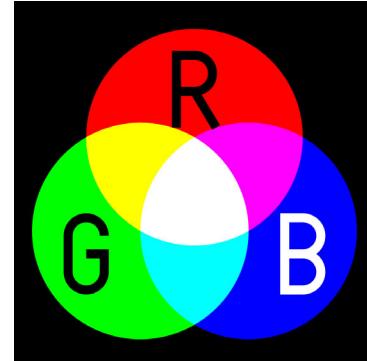


1	242	242	242	242	242	242	242	242	165	199	255	179	160	255	255	247	193	148	148	196	255	255	247	160	188	255	255						
1	242	242	242	242	242	242	242	242	163	173	250	240	148	208	199	154	190	212	173	173	250	256	256	256	212	240	250	250					
1	242	242	242	242	242	242	242	242	207	148	255	255	166	148	154	199	247	255	225	148	225	255	255	255	255	160	188	255	255				
1	242	242	242	242	242	242	242	242	230	148	196	160	153	148	208	255	255	255	255	173	173	255	255	255	255	212	148	173	173				
1	242	242	242	242	242	242	242	242	155	148	181	240	183	154	247	255	255	255	255	255	255	255	255	255	255	255	240	188	148	154			
1	242	242	242	242	242	242	242	242	186	158	255	255	247	154	199	255	255	255	255	255	255	255	255	255	255	255	160	240	188	148			
1	242	242	242	242	242	242	242	242	224	148	232	255	255	199	154	247	255	255	255	255	255	255	255	255	255	255	255	255	255	173	173		
1	242	242	183	148	207	242	242	242	242	242	160	193	255	255	247	154	199	255	247	199	154	150	154	255	255	255	255	255	255	225	225		
1	242	242	224	148	154	224	242	242	242	242	189	199	255	255	255	199	154	204	154	199	233	154	199	255	255	255	255	255	255	255	255		
1	242	242	172	148	166	235	242	242	242	224	148	252	255	255	212	149	148	183	247	255	255	199	154	247	255	255	255	255	255	255	255		
1	242	242	207	148	148	181	242	242	242	242	156	186	225	160	148	171	148	240	255	255	255	255	247	154	199	255	255	255	255	255	255	255	
1	242	242	242	154	173	148	200	242	242	242	200	149	148	179	240	255	160	186	255	255	255	255	255	199	154	247	255	255	255	255	255	255	255
1	242	242	242	195	160	199	148	217	242	242	242	235	148	196	255	255	255	255	255	255	255	255	255	255	255	255	247	154	189	199	154		
1	242	242	242	242	242	242	242	242	235	148	212	179	154	230	242	242	183	160	255	255	255	255	173	255	255	255	255	255	255	255	255	255	
1	242	242	242	183	156	247	188	186	230	242	242	148	248	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255		
1	242	242	242	242	242	242	242	242	166	148	232	242	148	248	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255		
1	242	242	242	242	242	242	242	242	177	173	255	225	148	242	212	148	240	255	255	255	255	255	173	163	160	148	188	240	212	148	240		
1	242	242	242	242	242	242	242	242	217	148	225	212	148	207	244	160	186	225	173	148	173	179	160	255	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	172	173	255	255	255	198	148	224	207	148	148	173	225	255	255	255	247	154	199	255	255	255	255	255	255
1	242	242	242	242	242	242	242	242	217	148	240	255	255	179	154	174	258	255	255	255	255	255	189	160	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	172	148	255	255	247	166	160	203	148	240	255	255	255	240	148	205	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	217	148	240	255	255	247	166	168	159	158	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
1	242	242	242	242	242	242	242	242	172	173	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	217	148	240	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	217	148	240	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	
1	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	242	

Color Space

- RGB

- ▶ RGB color images are created using 3 grayscale images
 - ▶ Red, green, and blue color channel values range from 0 - 255 due to 8-bit restrictions
 - ▶ RGB produces colors using additive color mixing



https://en.wikipedia.org/wiki/RGB_color_model#/media/File:AdditiveColor.svg

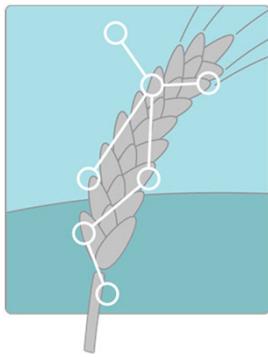
Image Formation

What is a pixel?



Image Formation

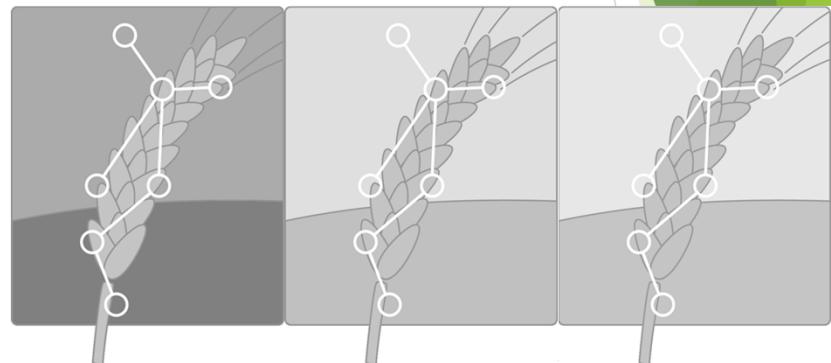
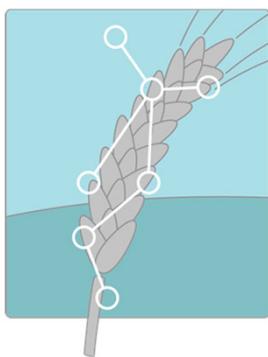
Color Space



- RGB - Red, Green, Blue color channels
- ▶ Uses 3 images to construct color
 - ▶ Each image is created with intensity values (grayscale)
 - ▶ Order for each image matters
 - ▶ Computer automatically assumed RGB order of images

Image Formation

Color Space



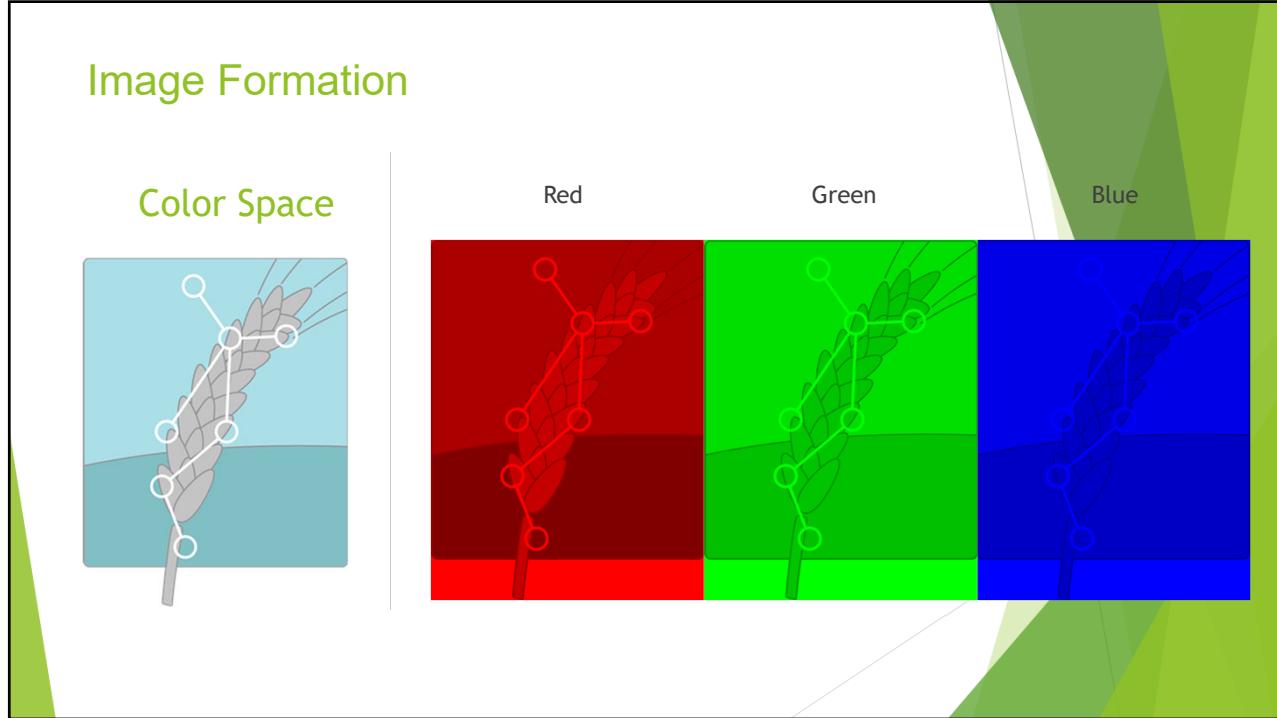
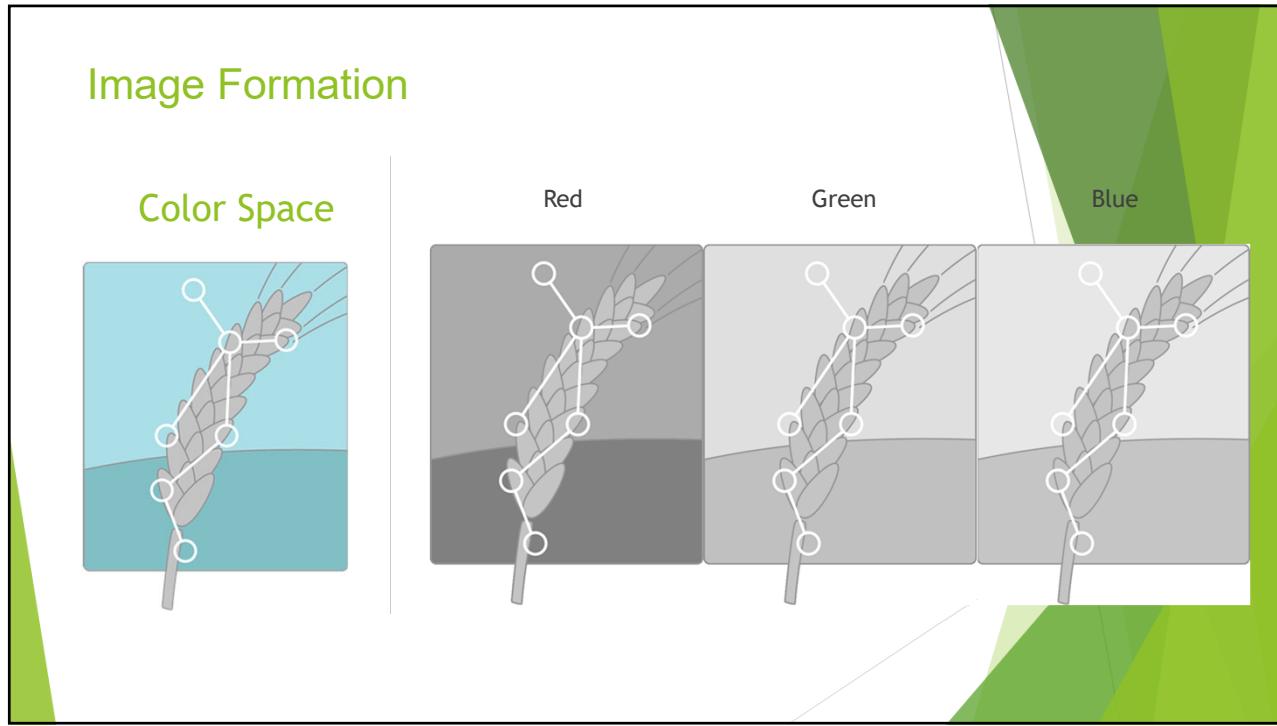
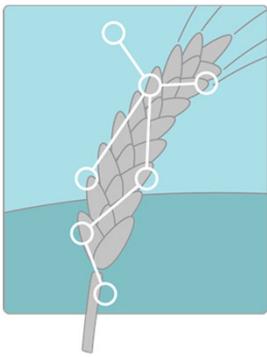


Image Formation

Color Space



Images combined in RGB order

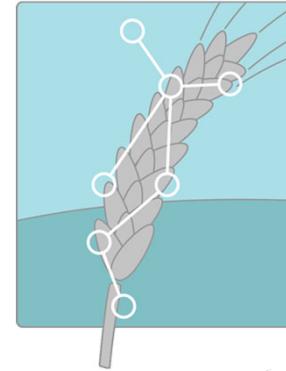
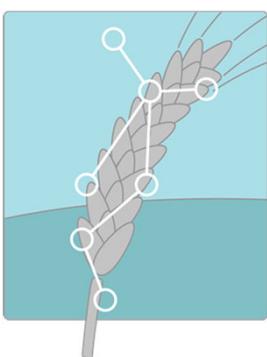
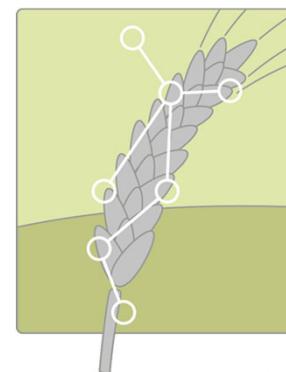


Image Formation

Color Space



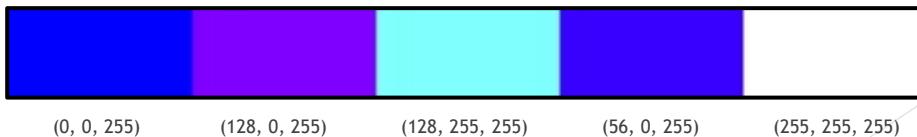
Images combined in GBR order



Color Space

- RGB

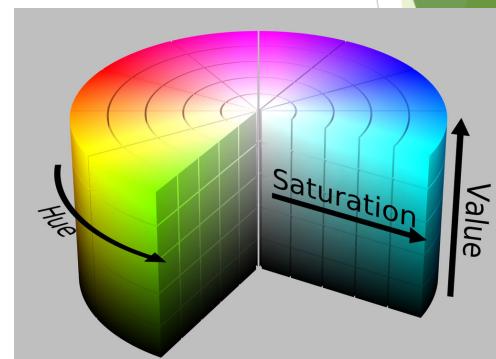
- ▶ Most commonly used color space
- ▶ OpenCV uses BGR instead of the traditional RGB order
- ▶ The “distance” between colors is very large and hard to group
- ▶ Below example shows the limitation of analyzing RGB color space:



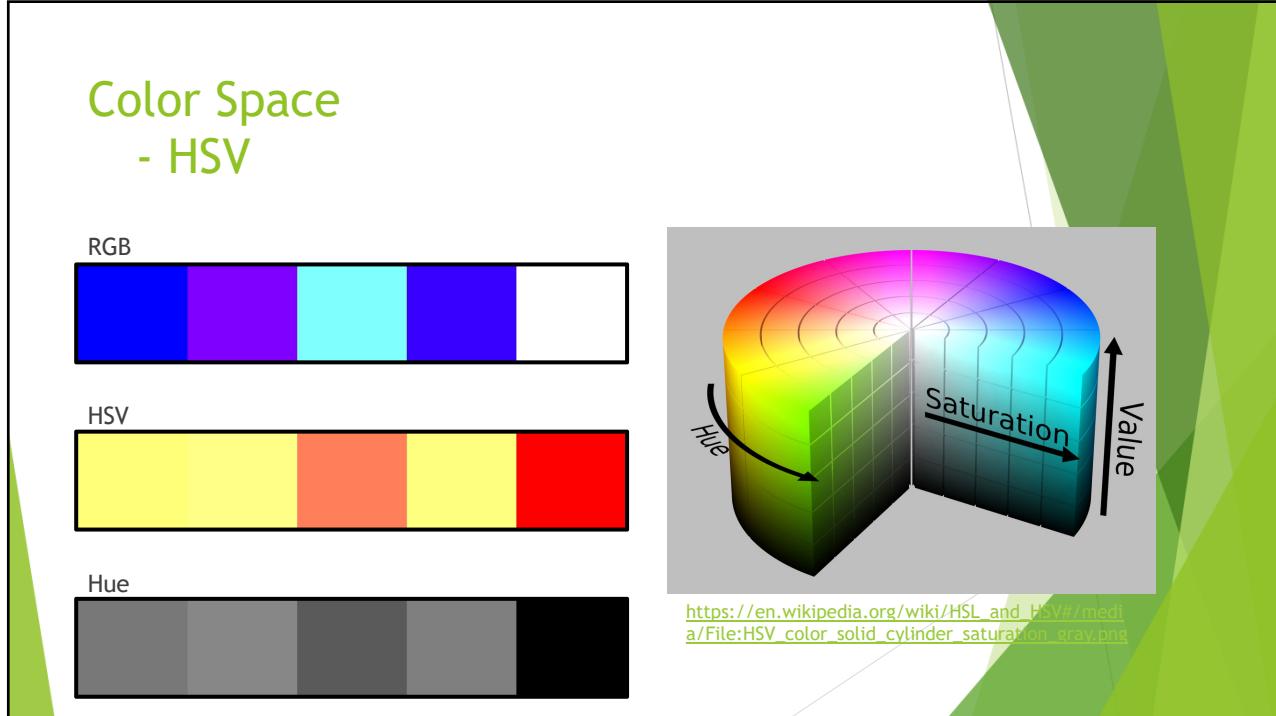
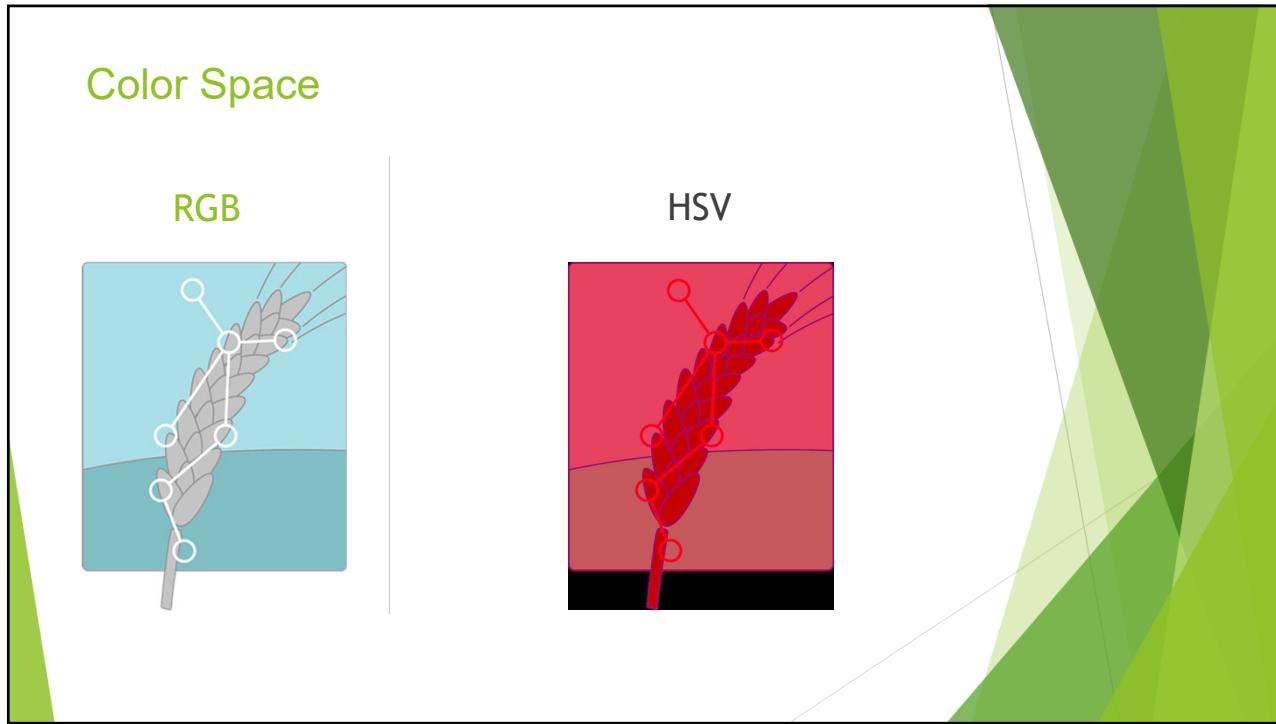
Color Space

- HSV

Real World	OpenCV (8-bits)
▶ H - Hue: 0 - 359 degrees	0 - 179
▶ S - Saturation: 0 - 100%	0 - 255
▶ V - Value: 0 - 100%	0 - 255
▶ Hue is the primary color information	
▶ Saturation represents the amount of gray	
▶ Value represents the intensity of the color	

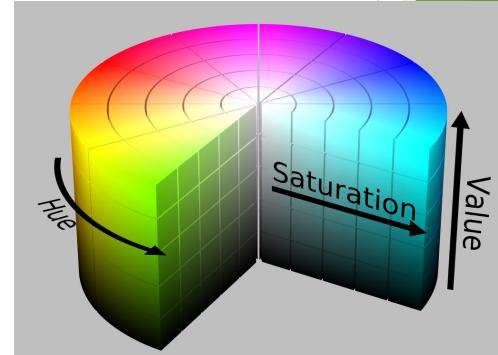
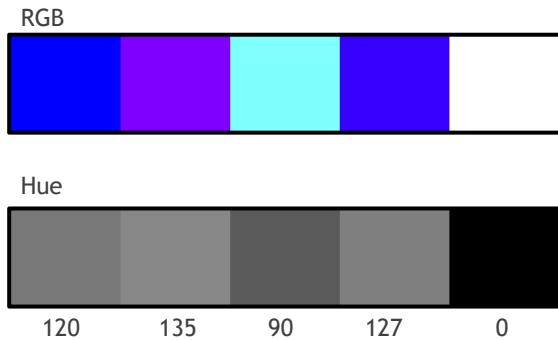


https://en.wikipedia.org/wiki/HSL_and_HSV#/media/File:HSV_color_solid_cylinder_saturation_gray.png



Color Space

- HSV



https://en.wikipedia.org/wiki/HSL_and_HSV#/media/File:HSV_color_solid_cylinder_saturation_gray.png

Image Thresholding

- Simple Thresholding

- If pixel value is greater than a threshold value, it is assigned one value (usually 255, white), else it is assigned to 0 (black).
- Image is transformed into a form of a binary image.

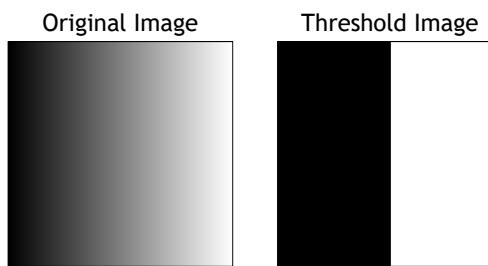


Image Thresholding

- **inRange Thresholding**

- ▶ If pixel value is in between the lower and upper threshold values, it is assigned 255 (white), else it is assigned the value 0 (black).
- ▶ Image is transformed into a form of a binary image.

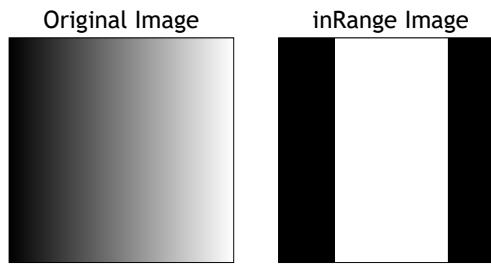


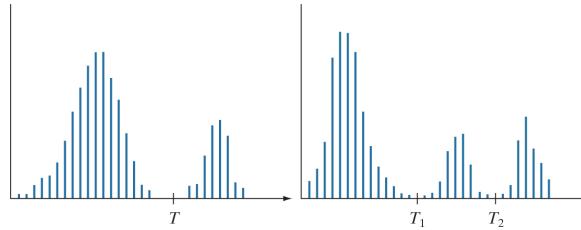
Image Thresholding

- ▶ When T is a constant applicable over an entire image, the process given in the equation is referred to as ***global thresholding***.
- ▶ When the value of T changes over an image, we use the term ***variable thresholding***. The terms local or regional thresholding are used sometimes to denote variable thresholding in which the value of T at any point (x, y) in an image ***depends on properties of a neighborhood of (x, y)*** (for example, the average intensity of the pixels in the neighborhood).
- ▶ If ***T depends on the spatial coordinates (x, y) themselves***, then variable thresholding is often referred to as ***dynamic or adaptive thresholding***.

Image Thresholding

a

b



- ▶ Suppose that the intensity histogram corresponds to an image, $f(x, y)$, composed of light objects on a dark background, in such a way that object and background pixels have intensity values grouped into two dominant modes. One obvious way to extract the objects from the background is to select a threshold, T , that separates these modes.

$$g(x, y) = \begin{cases} 1 & \text{if } f(x, y) > T \\ 0 & \text{if } f(x, y) \leq T \end{cases}$$

$$g(x, y) = \begin{cases} a & \text{if } f(x, y) > T_2 \\ b & \text{if } T_1 < f(x, y) \leq T_2 \\ c & \text{if } f(x, y) < T_1 \end{cases}$$

Image Thresholding - Adaptive Thresholding

- ▶ In simple thresholding, a global value is used as threshold value. But it may not be good in all the conditions where image has different lighting conditions in different areas. In that case, we use **adaptive thresholding** in which the algorithm calculates the threshold for a small regions of the image. So we get different thresholds for different regions of the same image and it gives us better results for images with varying illumination.
- ▶ Adaptive Method - Computationally decides how the thresholding value is calculated.
 - ▶ Mean - Threshold is determined by the mean of the neighborhood area.
 - ▶ Gaussian - Threshold is determined by the weighted sum of the neighborhood values. The weights are determined by a gaussian window.

Image Thresholding

- Adaptive Thresholding

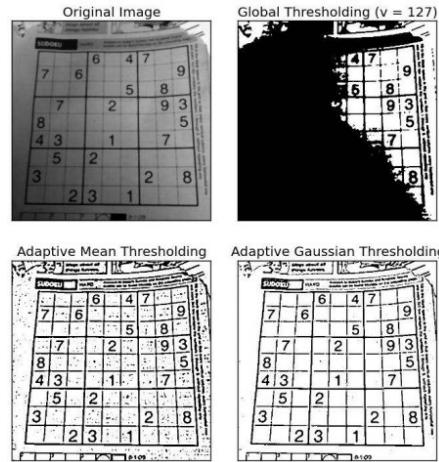


Image Thresholding

- Otsu Thresholding

- ▶ In global thresholding, we used an arbitrary value for threshold value. So, how can we know a value we selected is good or not?
- ▶ Answer is, *trial and error* method.
- ▶ But consider a bimodal image (*In simple words, bimodal image is an image whose histogram has two peaks*). Otsu binarization can automatically calculates a threshold value from image histogram for a bimodal image by *maximizing between-class variance* term.

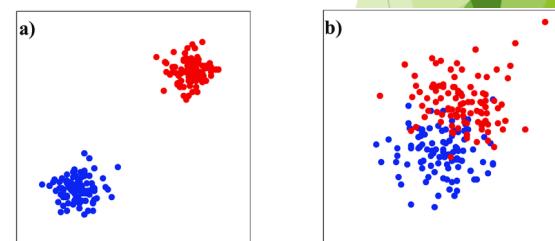


Image Thresholding

- Otsu Thresholding

- (a) Original image.
 - (b) Histogram (high peaks were clipped to highlight details in the lower values).
 - (c) Thresholding results using the basic global algorithm
 - (d) Result using Otsu's method.

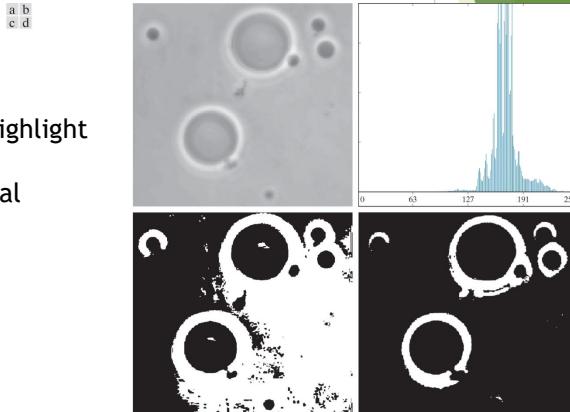
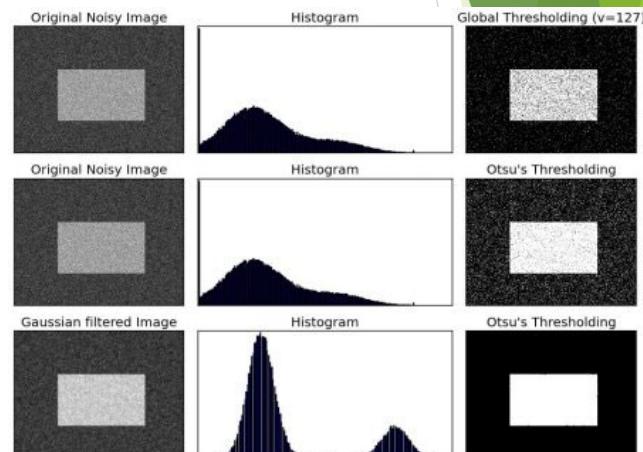


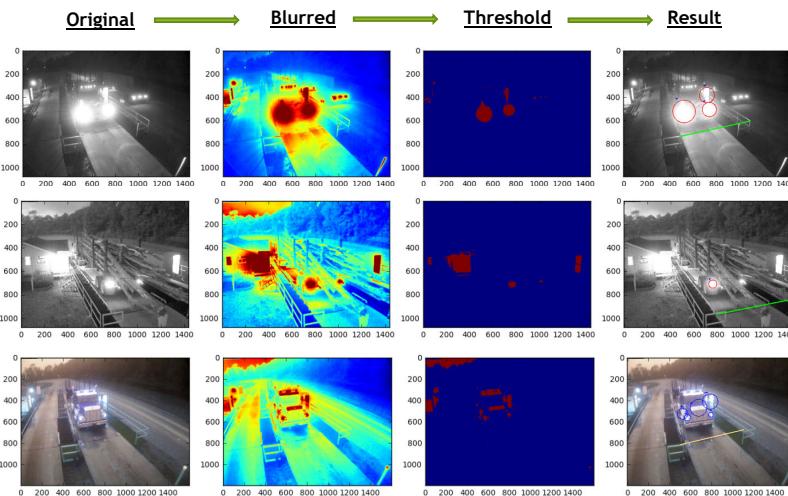
Image Thresholding - Otsu Thresholding

- ▶ Global thresholding
 - ▶ Otsu Thresholding
 - ▶ Otsu thresholding with Gaussian smoothing



Thresholding Applications

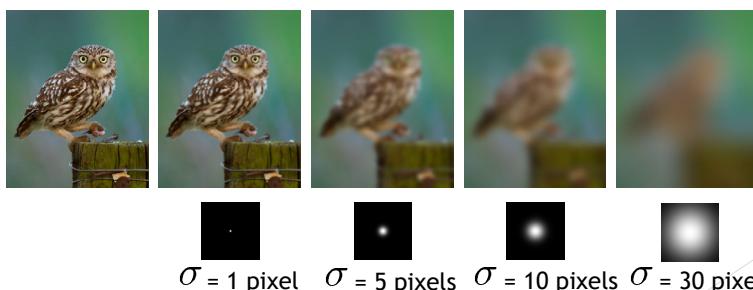
- Detect truck head/tail lights, and compare light's location with scale line
 - Apply different threshold values to detect head/tail lights at night and early morning



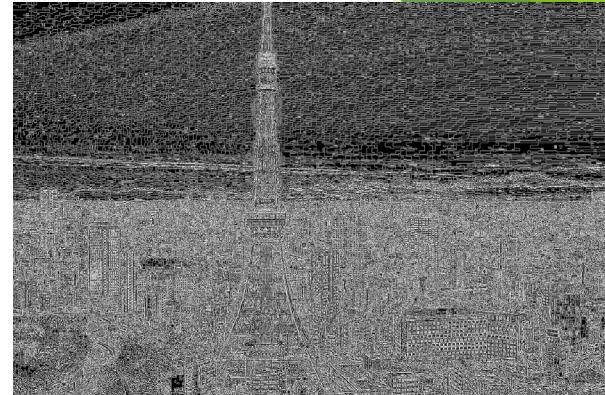
Accuracy: 98%

Blurring

- Why blur?
 - Smooths image to remove general noise from an image
 - Can help separate background and foreground features
 - Generally a good step before doing an initial analysis on an image

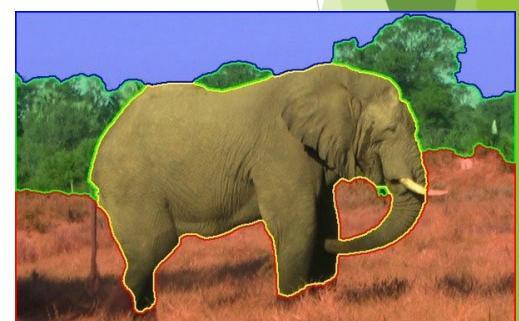


Spatial - LaPlacian



Feature Extraction

- ▶ After an image has been segmented into regions or their boundaries using segmentation methods, the resulting sets of segmented pixels usually have to be converted into a form suitable for further computer processing.
- ▶ Typically, the step after segmentation is feature extraction, which consists of feature detection and feature description.



Feature Extraction

- ▶ Feature detection refers to finding the features in an image, region, or boundary.
- ▶ Feature description assigns quantitative attributes to the detected features.
- ▶ For example, we might detect corners in a region boundary, and describe those corners by their orientation and location.
- ▶ Question: what is a feature in an image?



Image Features

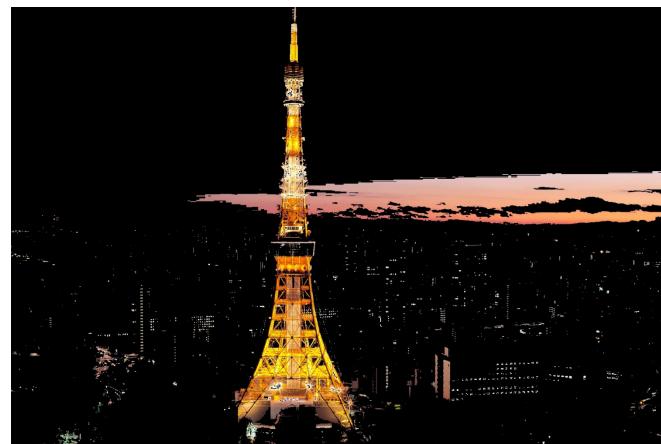
- ▶ In computer vision, image feature refers to two possible entities:
 - ▶ A global property of an image or part thereof, for instance the average gray level (global feature); or
 - ▶ A part of the image with some special properties, for instance a circle, a line, or a texture region in an intensity image, a planar surface in a range image (local feature)
- ▶ Formal definition of Image Features
- ▶ Image features are local, meaningful, detectable parts of the image



Purpose of Features

- ▶ Features by themselves are seldom generated for human consumption, except in applications such as interactive image processing.
- ▶ In fact, some feature extraction methods generate tens, hundreds, or even thousands of descriptor values that would appear meaningless if examined visually.
- ▶ Instead, feature description typically is used as a preprocessing step for higher-level tasks, such as image registration, object recognition for automated inspection, searching for patterns (e.g., individual faces and/or fingerprints) in image databases, and autonomous applications, such as robot and vehicle navigation.

Spatial - Background Removal



Spatial - “Foreground” Removal



Spatial - cv2.Laplacian after BG removal



Morphological Transformations

- ▶ Leverages kernels/structuring elements to perform simple operations
- ▶ Uses any kernel of $M \times N$ shape
- ▶ Generally uses odd numbered $M \times N$ shapes where $M = N$
- ▶ The values in each kernel should be binary (0 or 1)
- ▶ Two basic examples of morphological transformations
 - ▶ Erosion
 - ▶ “Removes” (chips away) non-zero pixels from non-zero pixel regions, based on the kernel shape
 - ▶ Dilating
 - ▶ “Expands” non-zero pixels, based on the kernel shape

Morphological Transformations

- Erosion

1. Establish a kernel and an “anchor” point (kernel center)
 2. Slide kernel across the image from top-left to bottom-right, centering the anchor point on each pixel
 3. At each pixel, evaluate the kernel overlaying the image
 - ▶ If at each non-zero kernel value, there is a non-zero pixel value on the overlaid image pixel, do nothing
 - ▶ Else if **any** non-zero kernel value does not have a corresponding non-pixel value on the overlaid image pixel, set the anchor point pixel of the image to 0
- ▶ Final result:
- ▶ Operation will only leave larger clusters of pixels and tend to remove smaller objects that could be associated with noise

Morphological Transformations

- Erosion

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

1	1	1	1	0	0
0	1	1	1	0	0
1	0	1	1	1	1
0	0	0	1	0	0
1	1	0	0	1	0

Final
Output

0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0
0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Morphological Transformations

- Erosion

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

1	1	1	0	0
0	1	1	1	0
1	0	1	1	1
0	0	0	1	0
1	1	0	0	1

Final
Output

0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0
0	0	0	0	0	0
0	0	0	0	0	0
0	0	0	0	0	0

Morphological Transformations

- Erosion

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

1	1	1	0	0
0	1	1	1	0
1	0	1	1	1
0	0	0	1	0
1	1	0	0	1

Final
Output

0	1	0	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0
0	0	0	0	0

Morphological Transformations

- Erosion

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

1	1	1	0	0
0	1	1	1	0
1	0	1	1	1
0	0	0	1	0
1	1	0	0	1

Final
Output

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Morphological Transformations

- Erosion

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

1	1	0	1	0
0	1	1	1	0
1	0	1	1	1
0	0	0	1	0
1	1	0	0	1

Final
Output

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Morphological Transformations

- Dilation

1. Establish a kernel and an “anchor” point (kernel center)
 2. Slide kernel across the image from top-left to bottom-right, centering the anchor point on each pixel
 3. At each pixel, evaluate the kernel overlaying the image
 - ▶ At each non-zero kernel value, set the corresponding overlaid image pixels to 1
- ▶ Final result:
- ▶ Operation will expand non-zero pixels into their surrounding regions, resulting in an image with larger clusters of non-zero pixels

Morphological Transformations

- Dilation

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Dilation
Operation

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Final
Output

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	0	1	0
0	0	0	0	0

Morphological Transformations

- Dilation

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Dilation
Operation

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Final
Output

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	0	1	0
0	0	0	0	0

Morphological Transformations

- Dilation

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Dilation
Operation

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Final
Output

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	0	1	0
0	0	0	0	0

Morphological Transformations

- Dilation

3x3 Kernel
with a cross pattern

0	1	0
1	1	1
0	1	0

Erosion
Operation

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0
0	0	0	0	0

Final
Output

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	0	1	0
0	0	0	0	0

Morphological Transformations

- Dilation

3x3 Kernel
with a cross pattern

Dilation
Operation

Final
Output

0	1	0
1	1	1
0	1	0

0	1	0	0	0
0	0	1	0	0
0	0	0	1	0
0	0	0	0	0

1	1	1	0	0
0	1	1	1	0
0	0	1	1	0
0	0	0	1	0
0	0	0	0	0

Dilation



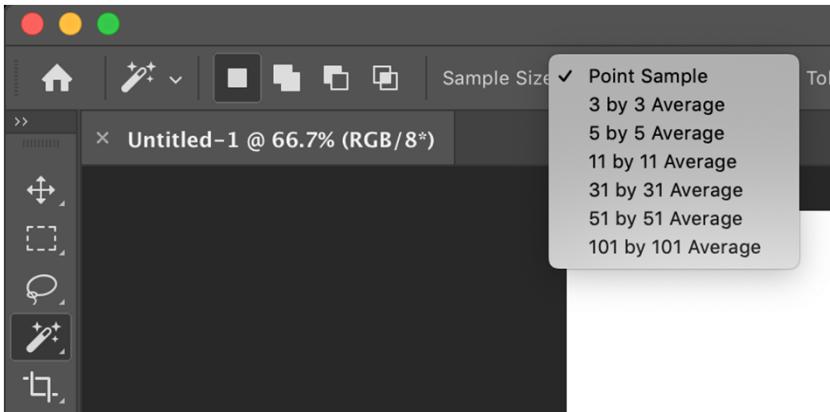
(a) Initialization.



(b) Result.



Show Photoshop



Masks

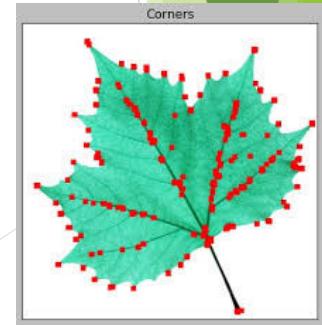
- ▶ Binary black and white images
- ▶ Bitwise operations, such as `bitwise_and` can be used to obtain parts of the original image
- ▶ Most segmentation operations use masks instead of the original image



Harris Corner Detection

- ▶ Corners are regions in the image with large variation in intensity in all the directions.
- ▶ One early attempt to find these corners was done by Chris Harris & Mike Stephens in their paper A Combined Corner and Edge Detector in 1988, so now it is called Harris Corner Detector.
- ▶ The approach is to maximize

$$C(x, y) = \sum_s \sum_t w(s, t)[f(s + x, t + y) - f(s, t)]^2$$



Harris Corner Detection

$$C(x, y) = \sum_s \sum_t w(s, t)[f(s + x, t + y) - f(s, t)]^2$$

- ▶ The weighting function $w(x, y)$ used in the HS detector generally has one of two forms:
 - ▶ It is 1 inside the patch and 0 elsewhere (i.e., it has the shape of a box lowpass filter kernel), or
 - ▶ It is an exponential function of the form
- $$w(s, t) = e^{-(s^2+t^2)/2\sigma^2}$$
- ▶ The box is used when computational speed is paramount and the noise level is low. The exponential form is used when data smoothing is important.

Harris Corner Detection

$$C(x, y) = \sum_s \sum_t w(s, t) [f(s + x, t + y) - f(s, t)]^2$$

$$C(x, y) = [x \ y] M \begin{bmatrix} x \\ y \end{bmatrix}$$

Where,

$$M = \sum_s \sum_t w(s, t) \begin{bmatrix} f_x^2 & f_x f_y \\ f_x f_y & f_y^2 \end{bmatrix}$$

Harris Corner Detection

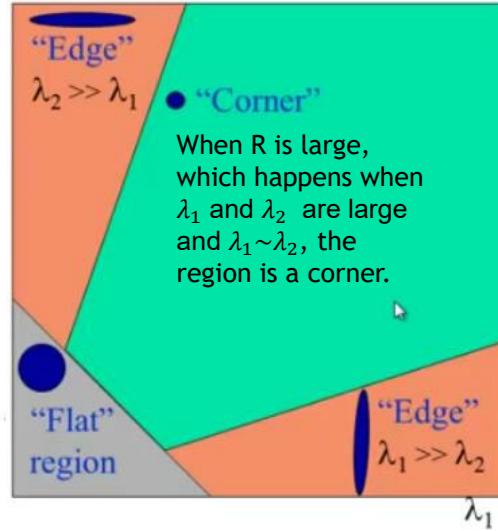
- ▶ $R = \det(M) - k(\text{trace}(M))2$
- ▶ $\det(M) = \lambda_1 \lambda_2$
- ▶ $\text{trace}(M) = \lambda_1 + \lambda_2$
- ▶ λ_1 and λ_2 are eigenvalues of M .
- ▶ Constant k is determined empirically, and its range of values depends on the implementation; In general, the smaller it is, the more likely to find corners.

Harris Corner Detector

$$R = \det(M) - k(\text{trace}(M))^2$$

When $R < 0$, which happens when $\lambda_1 \gg \lambda_2$ or vice versa, the region is edge.

When $|R|$ is small, which happens when λ_1 and λ_2 are small, the region is flat.

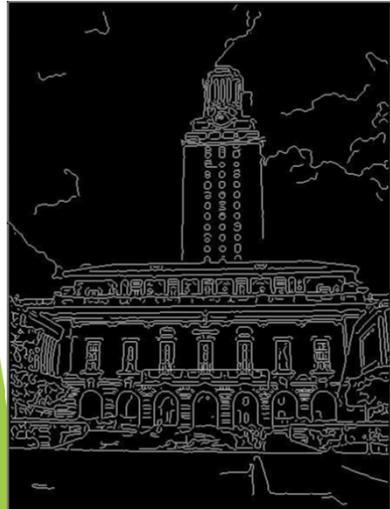


Hough Transform

- ▶ A popular algorithm to detect lines and simple curves.

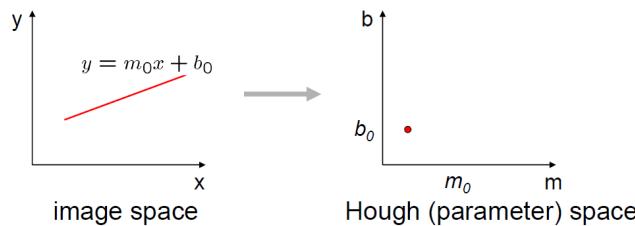


Hough Transform



- ▶ Hough Transform is a voting technique that answers these questions:
 - Given points that belong to a line, what is the line?
 - How many lines are there?
 - Which points belong to which lines?
- ▶ General Methodology:
 - Record vote for each possible line on which each edge point lies.
 - Look for lines that get many votes.

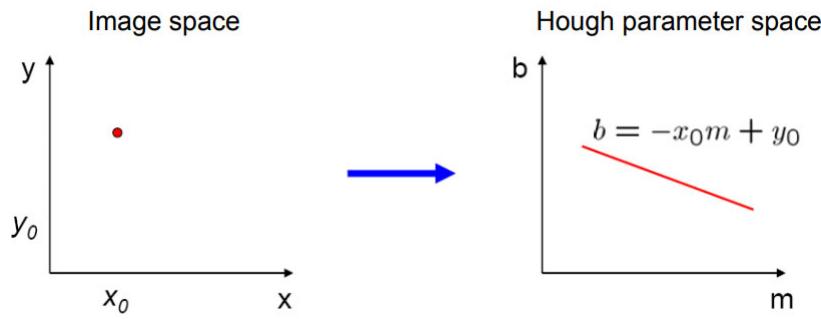
Hough Transform



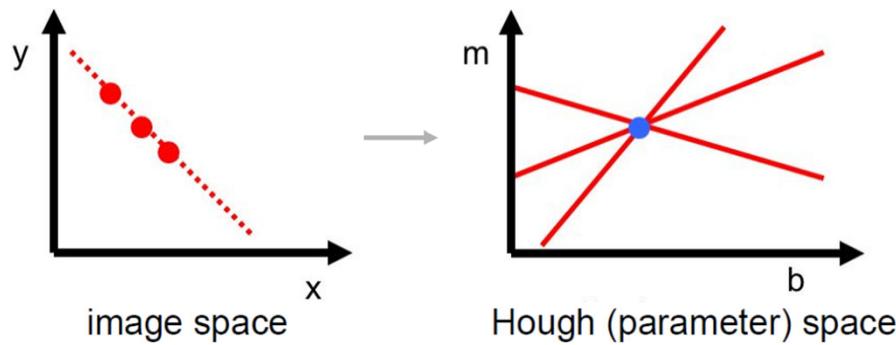
Connection between image (x,y) and Hough (m,b) spaces

- A line in the image corresponds to a point in Hough space
- To go from image space to Hough space:
 - given a set of points (x,y) , find all (m,b) such that $y = mx + b$

Hough Transform



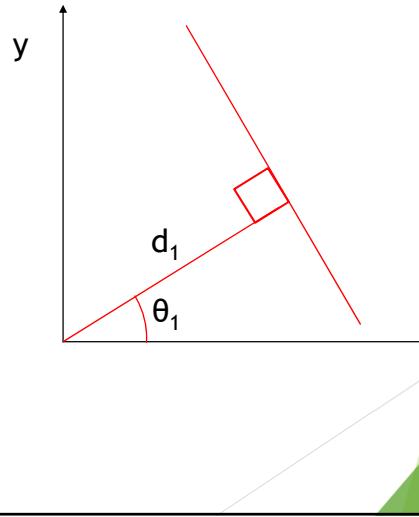
Hough Transform



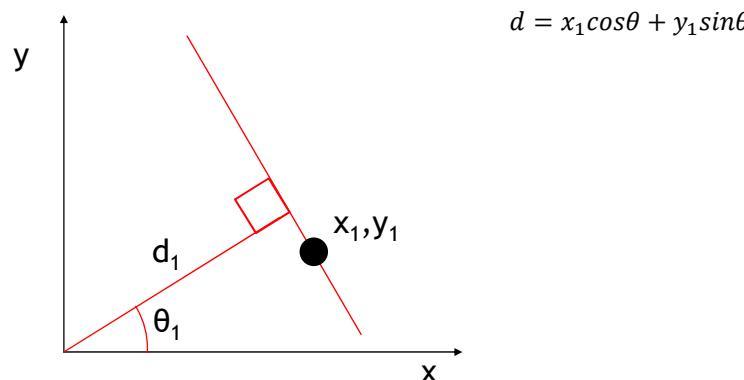
P.V.C. Hough, *Machine Analysis of Bubble Chamber Pictures*, Proc. Int. Conf. High Energy Accelerators and Instrumentation, 1959

The Hough Transform

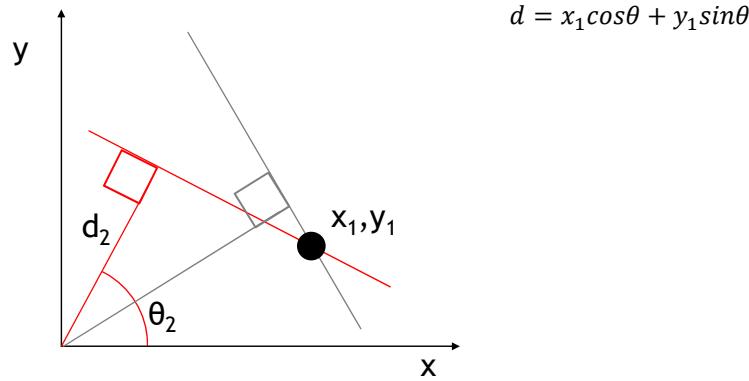
- ▶ Use polar representation to avoid infinite slope



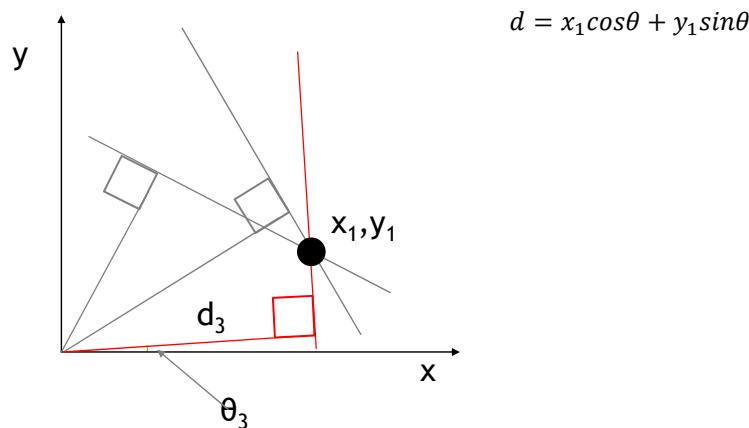
The Hough Transform



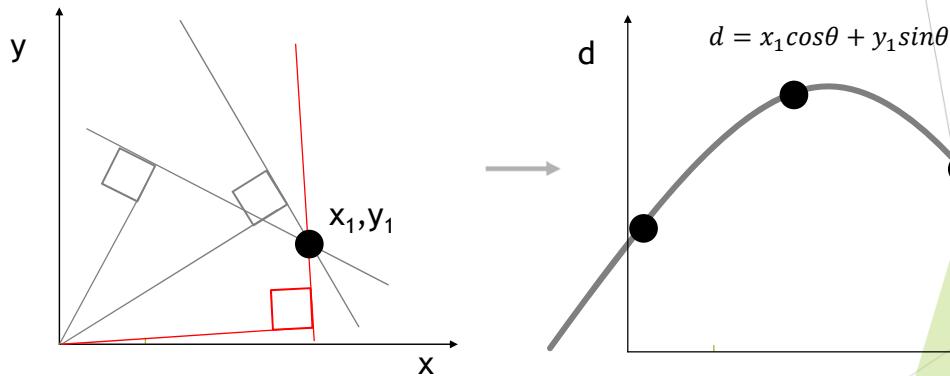
The Hough Transform



The Hough Transform



The Hough Transform



Hough Transform

Using the polar parameterization:

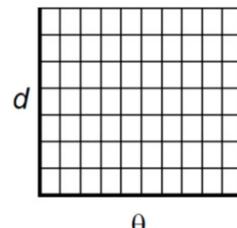
$$x \cos \theta - y \sin \theta = d$$

Basic Hough transform algorithm

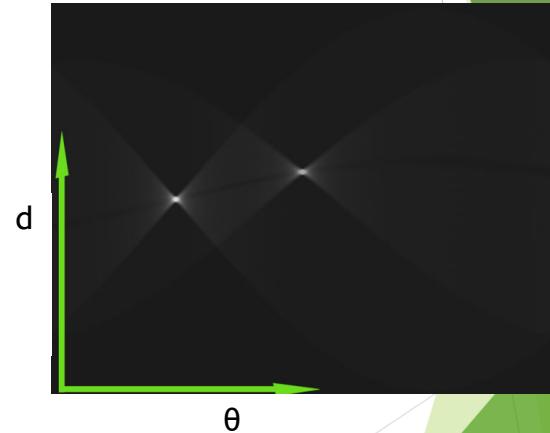
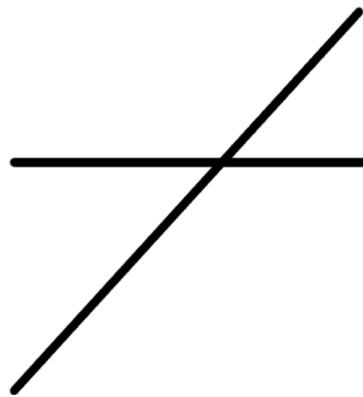
1. Initialize $H[d, \theta] = 0$
2. for each edge point $I[x, y]$ in the image


```
for θ = [θmin to θmax] // some quantization
        d = x cos θ - y sin θ
        H[d, θ] += 1
```
3. Find the value(s) of (d, θ) where $H[d, \theta]$ is maximum
4. The detected line in the image is given by $d = x \cos \theta - y \sin \theta$

H : accumulator array (votes)



The Hough Transform (HT)



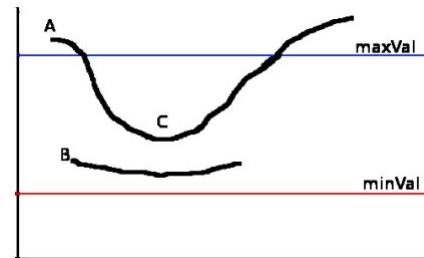
Contours and Edge Detection - findContours

- ▶ Contours are a set of edges that are enclosed
- ▶ findContours produces coordinate values for each set of contours
- ▶ Useful in obtaining features and manipulating results obtained from other edge detection algorithms

Contours and Edge Detection

- Canny

- ▶ Uses the sobel x, sobel y algorithms to obtain edges in both directions
- ▶ Determines the center of each edge and transforms the edges into a single pixel wide line
- ▶ Uses hysteresis thresholding, on the gradient intensity, to obtain the final valid edges
 - ▶ Any line below the minVal is dropped
 - ▶ Any line above maxVal is kept
 - ▶ Any line between minVal and maxVal is separately evaluated



Contours and Edge Detection

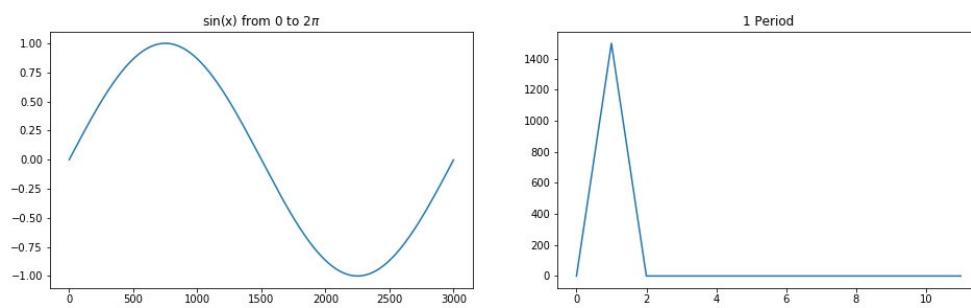
- boundingRect

- ▶ Finds the straight edge, unrotated, rectangle that bounds the contours
- ▶ Obtains:
 - ▶ x, y coordinates of the top left corner of the box
 - ▶ height, width of the rectangle
- ▶ Important to note:
 - ▶ OpenCV lists the coordinates as x, y
 - ▶ NumPy lists the coordinates as y, x (row, column)

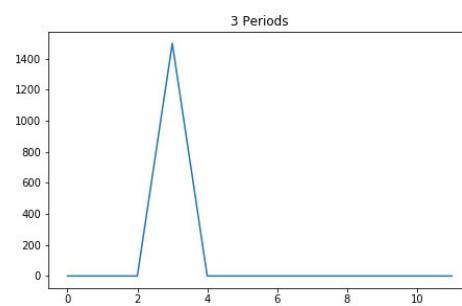
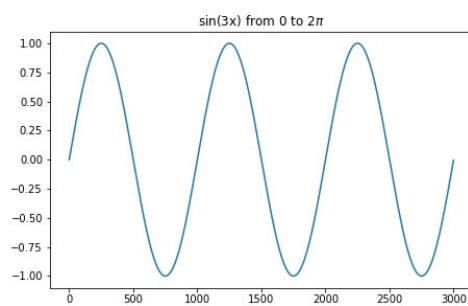
Cropping

- ▶ Cropping can be achieved through NumPy.
- ▶ `Image[yi:yf, xi:xf]`
 - ▶ For channel specific: `Image[yi:yf, xi:xf, channel]`
- ▶ Things to remember:
 - ▶ NumPy uses y, x format
 - ▶ When subsetting, NumPy does [i, f), so i is inclusive, but f is exclusive

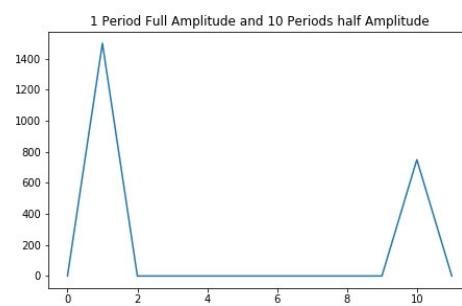
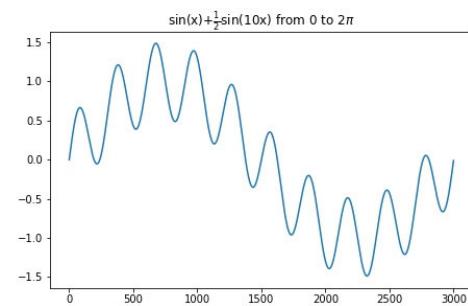
Frequency Domain - Fourier Transform



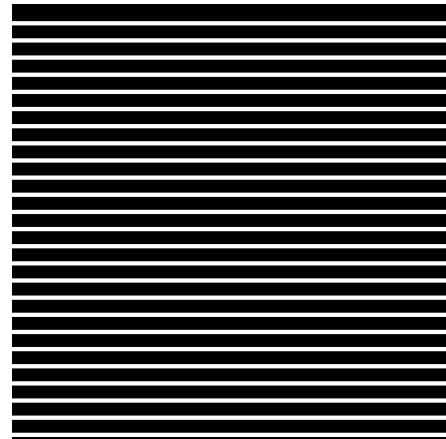
Frequency Domain - Fourier Transform



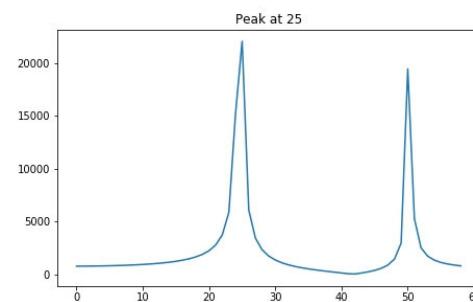
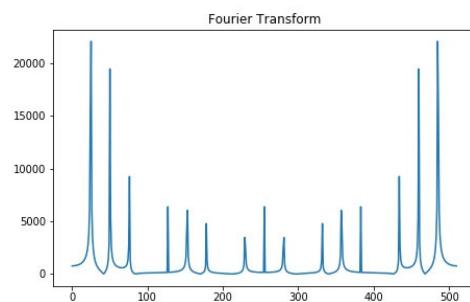
Frequency Domain - Fourier Transform



Frequency Domain - Fourier Transform



Frequency Domain - Fourier Transform

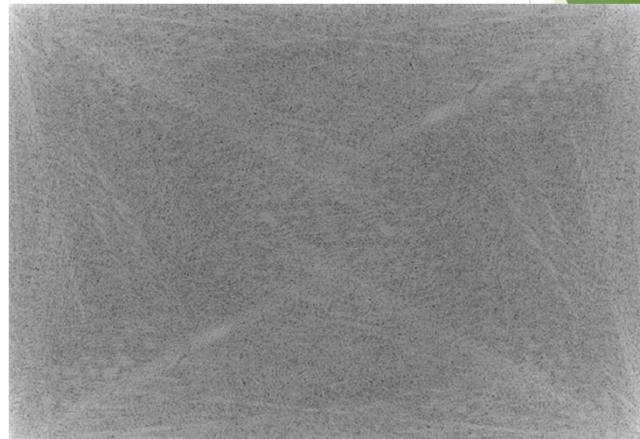


Fourier Transforms

FFT in Two Dimensions



- ▶ Output plots frequency_x, frequency_y, power (amplitude)



Fourier Transforms

FFT in Two Dimensions

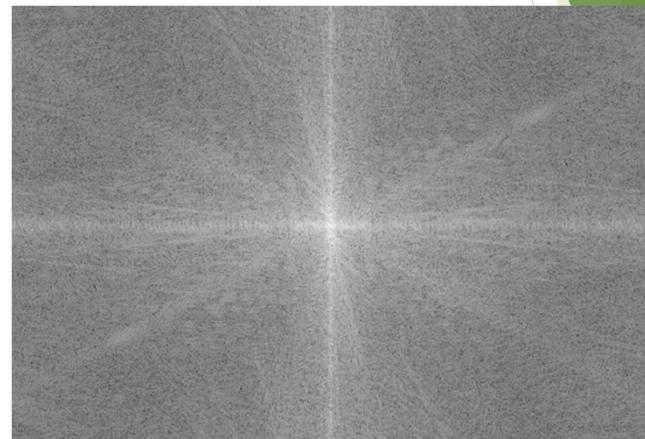
- ▶ Shifting centers the low frequencies to the center
- ▶ FFTs can be altered to generate high-pass and low-pass filters
- ▶ Manipulations in the Fourier domain can be approximately accomplished in the spatial domain and vice versa

Fourier Transforms

FFT in Two Dimensions



► Shifted FFT image

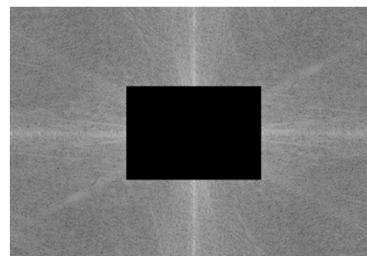


Fourier Transforms

FFT in Two Dimensions



► High Pass Filter

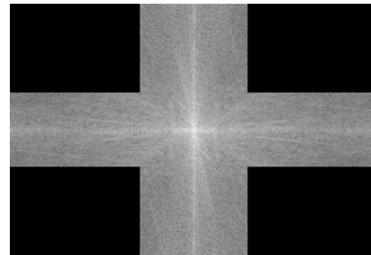


Fourier Transforms

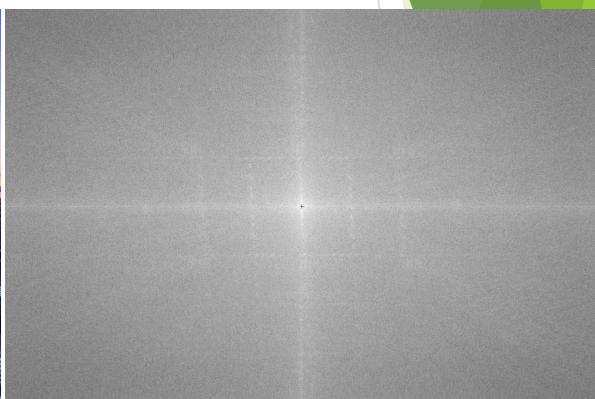
FFT in Two Dimensions



► Low Pass Filter



Spatial Vs Frequency



<https://www.burgessyachts.com/en/locations/tokyo/>

Frequency



Spatial - Kernel

